

Scalable approach for Learning Word Representations

Name: Zhang, Zijian
Matrikelnummer: 3184680


First Examiner: Prof.Dr. Avishek Anand
Second Examiner: Prof. Dr. techn. Dipl.-Ing. Wolfgang Nejdl
Advisor: Prof. Dr. Avishek Anand

Content


- Background Knowledge Introductions
- Structure of our Approach
- Original Corpus Division
- Merging Approaches
 - Concatenation+PCA
 - Manifold Alignment + Vector Averaging
- Experiments Setup
- Performance and Timing
- Conclusion and Future Works

Background Knowledge Introduction


Word Vector Representation



Corpus



Representation model



(Local) Euclidean Vector Space

Background Knowledge Introduction

Why Word2Vec:

Word representations:

used to be processed in computer:

One-hot / hash tokenizing (-lose of semantic information)

LSA: directly SVD on word-doc matrix (-linear model)

Bag-of-Word: meaning defined by set of labels (-low semantic accuracy, supervised)

Words co-occurrence matrix: count of word appears together (-sparse)

Therefore Word2Vec: a predictive model

preserves semantic information

dense representations for larger corpus

non-linear dimension reduction, relationship-preserving and robust

unsupervised way

a neural approach is more scalable

popular in downstream applications

Background Knowledge Introduction

Challenge of Word2Vec:

larger corpus: more **expressive**, longer **time**

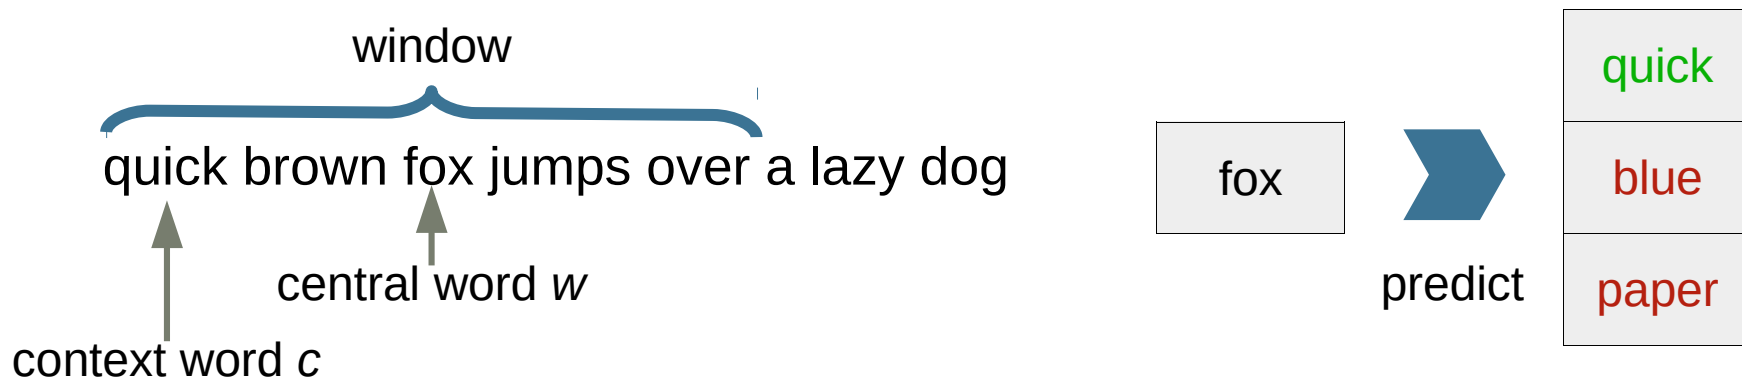
On 14GB English Wikidump: \approx **36 hours** !

(dim: 500 Vocabulary: 300k)

Not **scalable** w.r.t. large corpus

Background Knowledge Introduction


Skip Gram Negative Sampling




$$J(\theta) = \sum_{w \in D} \sum_{c \in I(w)} \underbrace{\left\{ \ln[p((w, c) | \theta)] \right\}}_{\text{Skip Gram}} + E_{\bar{c} \in V} \underbrace{\ln[1 - p((w, \bar{c}) | \theta)]}_{\text{Negative Sampling}}$$

Background Knowledge Introduction


Properties of Word2Vec [1]



$\cos(\theta) \sim \text{similarity}$



categories \Rightarrow clusters



analogy \Rightarrow parallelization

[1]: Mikolov T, Yih W, Zweig G.
Linguistic regularities in continuous space word representations
Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies. 2013: 746-751.

Background Knowledge Introduction

Recently work:

Optimizing SGD schema (online, synchronized) :

HogWild!, BLAS-3, optimizing cache usage, column-wise distribution

Adaptive SGD :

SGD-Momentum, AdaGrad, RMSProp

Heterogeneous combination:

Combine vectors from Word2Vec, WordNet, GloVe, ConceptNet etc.


Incremental Word2Vec training

Background Knowledge Introduction

Our contribution:


TBF

Structure of our Approach



Structure of our Approach

Research Questions:



- 1) How do we sample from the original corpus?
- 2) How do we combine sub-models?

Original Corpus Division

Criterion of Division: “Mimicking” the Original Term Distribution

objective function:

$$J(\theta) = \sum_{w \in D} \sum_{c \in l(w)} \{ \ln[p((w, c) | \theta)] + E_{\bar{c} \in V} \ln[1 - p((w, \bar{c}) | \theta)] \}$$

local optimum:

$$\frac{\partial J(\theta)}{\partial \theta} = 0$$

final result ^[1]:

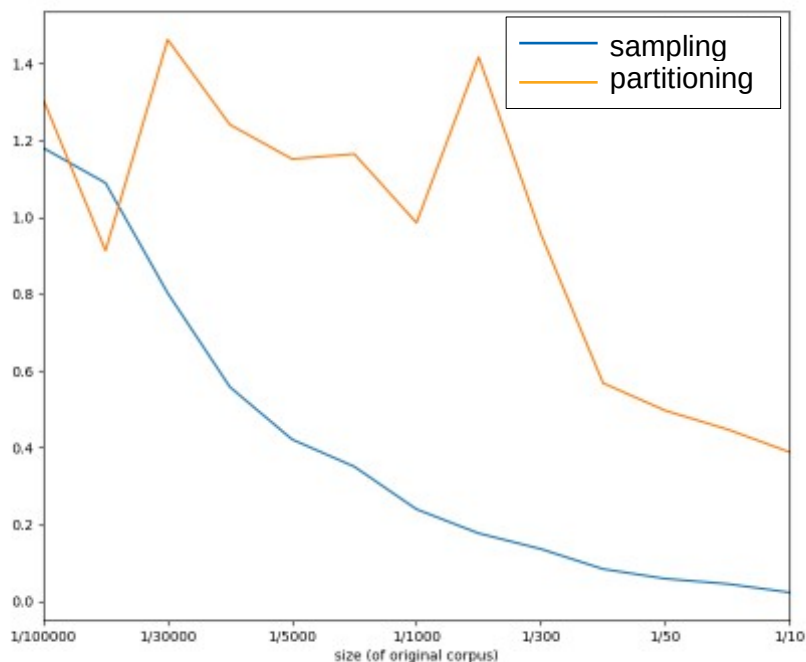
$$w c = \log \frac{N \#(w, c)}{\#(w) \#(c)} - \log k$$

Depends on distribution of words/central-contextual pairs only

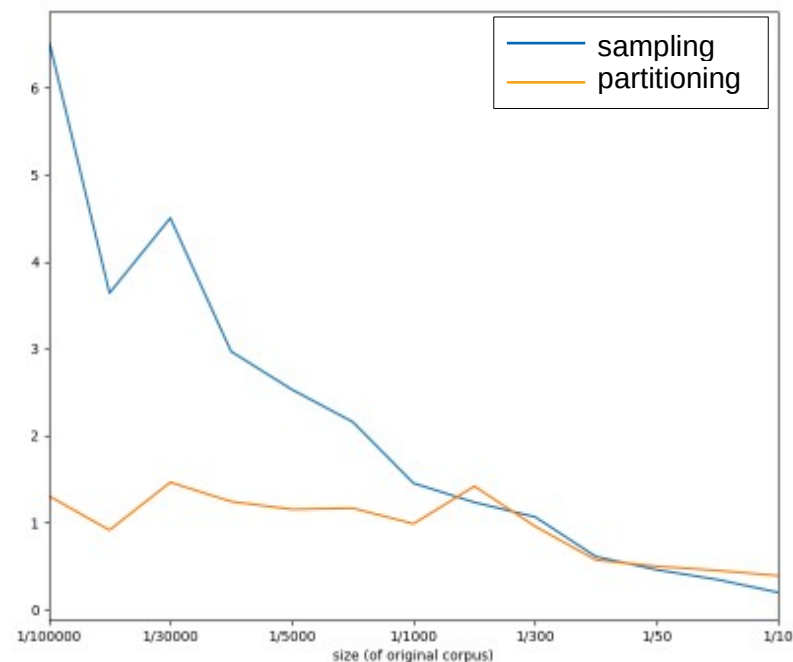
[1]: Levy O, Goldberg Y.
Neural word embedding as implicit matrix factorization
Advances in neural information processing systems. 2014: 2177-2185.

Original Corpus Division

Distribution preservation through uniform sampling



Word distribution



Central contextual words pair distribution

Merging Approaches


Problem so-far:



?



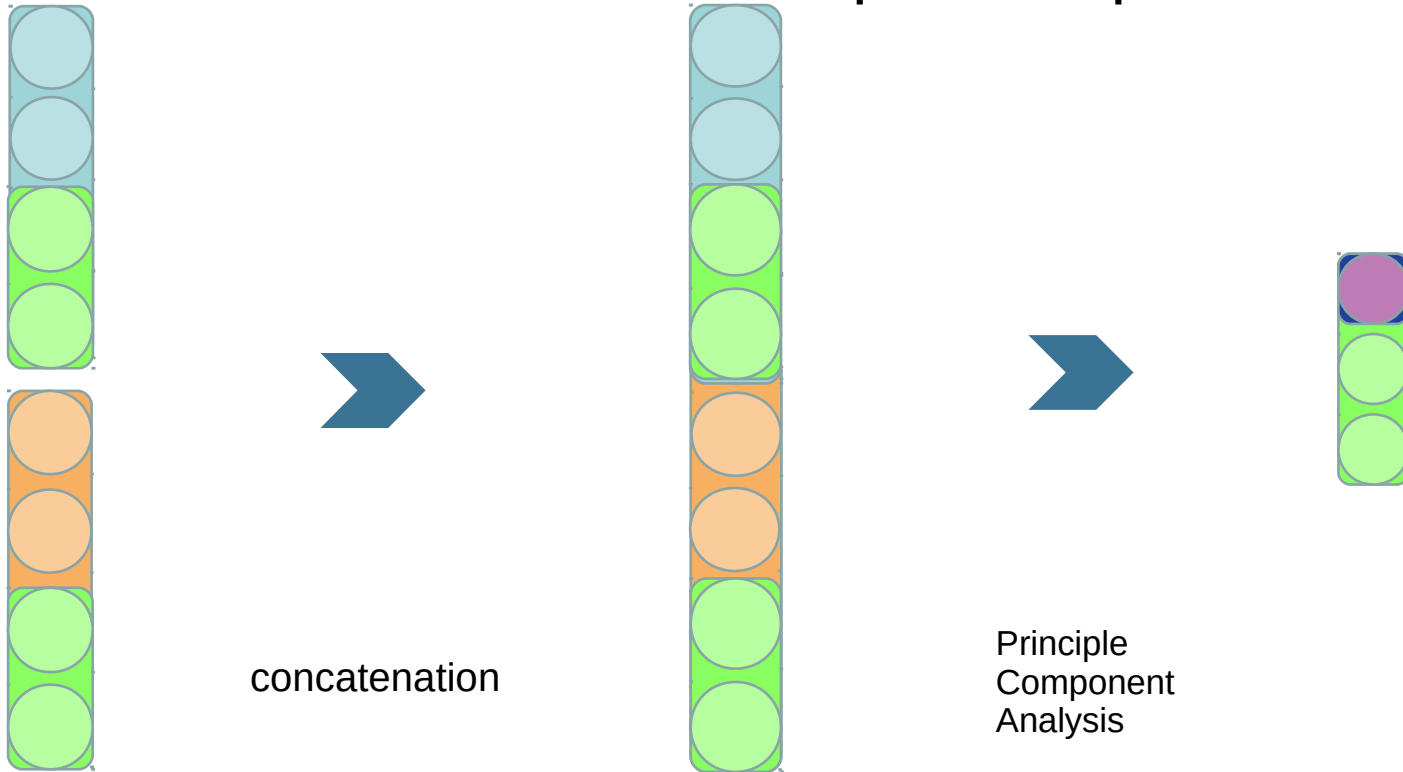
identify the
correlating parts



combine the correlation,
canceling noise by merging


Merging Approaches

1st Approach: Concatenation + Principle Component Analysis



Merging Approaches

Some Comments on PCA [1]



Law of cosines:

$$\|d\|^2 = \|a\|^2 + \|b\|^2 - 2\|a\|\|b\|\cos \theta$$

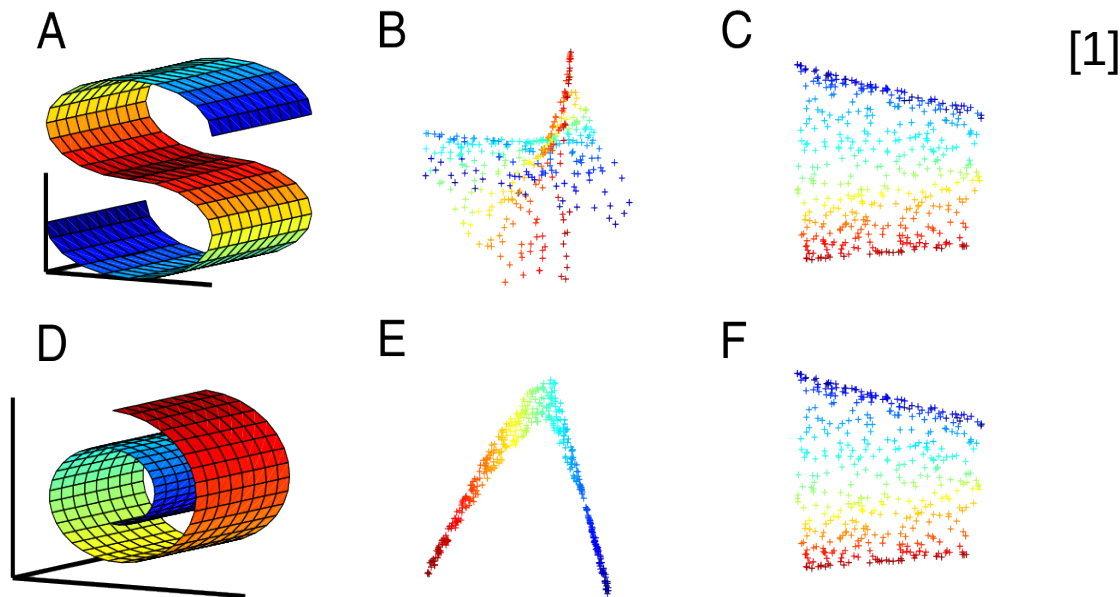
$$\|a\|^2 = \|b\|^2 = 1$$

$$\cos \theta = \frac{2 - \|d\|^2}{2}$$

[1]: Korenius T, Laurikkala J, Juhola M.
On principal component analysis, cosine and Euclidean measures in information retrieval
Information Sciences, 2007, 177(22): 4893-4905.

Merging Approaches


2nd Approach: Manifold Alignment + Vector Averaging



[1]: Ham, Ji Hun, Daniel D. Lee, and Lawrence K. Saul.
Learning high dimensional correspondences from low dimensional manifolds

Merging Approaches

2nd Approach: Manifold Alignment [1] + Vector Averaging



Find a low-rank correlation of Y

$$X_{\text{final}}, Y_{\text{final}} = (1-\mu)(\text{low-rank reconstruction}) + \mu(\text{inter-manifold correlation})$$

Find a low-rank correlation of Y

[1]: Boucher, Thomas, et al.
Aligning Mixed Manifolds
AAAI. 2015.

[2]: Ham J, Lee D D, Saul L K.
Semisupervised alignment of manifolds
AISTATS. 2005: 120-127.

Experiments Setup

Data Set:

- *Original corpus*: English Wikipedia dump file 2016 (14G)
- *Cleaner*: Wikipedia LineSentences cleaner in Gensim
- *Sample/Partition size*: 1%, 10% of the original corpus

Baseline Model^[1]:

- *Vocabulary*: 300k words
- *Dimension (size of hidden layer)*: 500
- *Training time*: 36 hours

Benchmarks:

- *Word Similarity*: MEN, RW, RG65, WS353, SimLex999, Mturk
- *Analogy*: Google, SemEval 2012 #2, MSR
- *Categorization*: AP, BLESS, Battig

[1]: Jastrzebski, Stanisław, Damian Leśniak, and Wojciech Marian Czarnecki
How to evaluate word embeddings? On importance of data efficiency and simple supervised tasks

Performance and Timing

Effect of Dividing (sampling/partitioning):

Approach	AP	Battig	MEN	RG65	RW	WS353	Google	SemEval2012_2
Partition(1/10,10)	0.582	0.420	0.719	0.725	0.232	0.596	0.615	0.162
Random(1/10,10)	0.604	0.435	0.740	0.757	0.209	0.637	0.654	0.188
Partition(1/100,100)	0.505	0.358	0.622	0.723	0.203	0.516	0.467	0.129
Random(1/100,100)	0.517	0.381	0.644	0.746	0.257	0.542	0.487	0.138
Baseline	0.595	0.434	0.736	0.757	0.299	0.611	0.661	0.181

Table 2: Comparison between Random(. , .) and Partition sampling approaches at different granularities. PCA is used as the merging method.

Conclusions:

Sampling performs overall better than **partitioning**

Larger division size makes the performance better, but less scalability

Surprisingly, some of our approaches even outperform the **baseline**!

Performance and Timing

Effect of Merging (concatenation/PCA/LRA):

Approach	AP	Battig	MEN	RG65	RW	WS353	Google	SemEval2012_2
Random(1/10,10) + Concat	0.595	0.429	0.742	0.765	0.268	0.611	0.645	0.187
Random(1/10,10) + LRA	0.560	0.327	0.685	0.785	0.264	0.606	0.661	0.158
Random(1/10,10) + PCA	0.604	0.435	0.740	0.757	0.209	0.637	0.654	0.188
Random(1/100,100) + Concat	0.525	0.377	0.643	0.746	0.257	0.542	0.487	0.138
Random(1/100,100) + LRA	–	–	–	–	–	–	–	–
Random(1/100,100) + PCA	0.517	0.381	0.644	0.746	0.257	0.542	0.487	0.138
Baseline	0.595	0.434	0.736	0.757	0.299	0.611	0.661	0.181

Table 3: Evaluation results for merging

Conclusions:



PCA performs overall the best

LRA and Vector average is the slowest

To see the information loss in the PCA, we added the concatenation as a single line, it performs surprisingly robust, despite its higher dimensionality

Performance and Timing

Effect of sample count (**10%** of original corpus):



Conclusion:

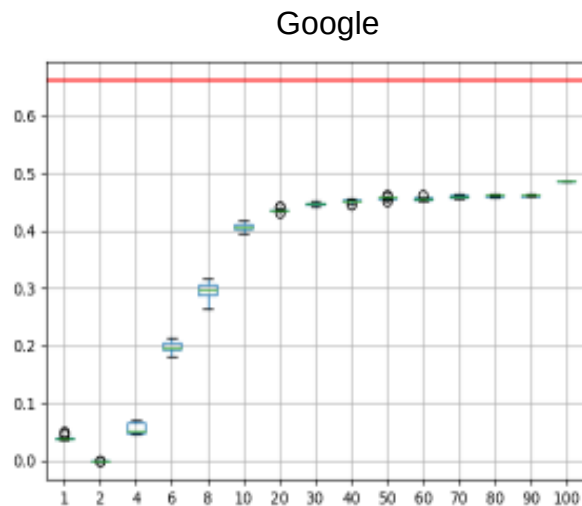
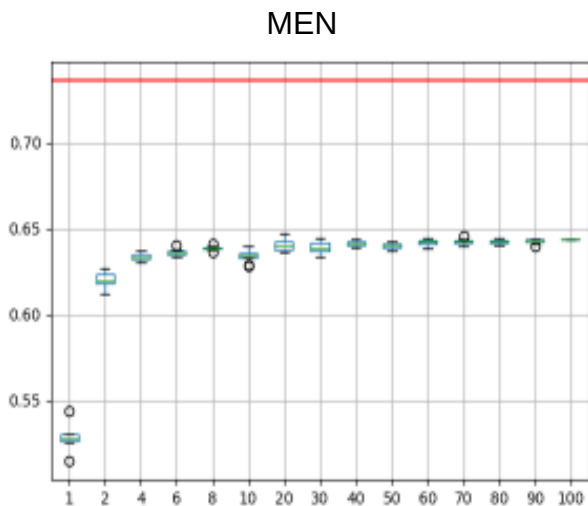
Larger count of sub-models doesn't help.

8~10 for **10%** approach, **20~30** for **1%** approach is enough.

After that it goes into flat road.

Performance and Timing

Effect of sample count (1% of original corpus):



Conclusion:

Larger count of sub-models doesn't help.

8~10 for 10% approach, 20~30 for 1% approach is enough.

After that it goes into flat road.

Performance and Timing

Timing:

Subject name	Time Elapse (s)
10% partitioning PCA	10932
1% partitioning PCA	2093
10% sampling PCA	12852
1% sampling PCA	16293
10% sampling LRA	27385
Baseline	129015

Conclusions:

m-part division leads to m-folds speed up

Total training time: from 36 hours @ baseline to
3.6 hours @ 10% and
< 1 hour @ 1% approach

Time of merging is in minutes

Conclusion and Future Works

Conclusions:

- Sentence-wise sampling to divide, PCA to merge
- M-fold speed up by M-splitting, sometimes even outperform the baseline.

Future works:

- Works on the go: dealing with missing-word problem between sub-models

Thank you for your attention

Questions and comments are highly appreciated

Zhang, Zijian
Hannover
08. 03. 2018