

Scalable approach for Learning Word Representations

Name: Zhang, Zijian
Matrikelnummer: 3184680

First Examiner: Prof.Dr. Avishek Anand
Second Examiner: Prof. Dr. techn. Dipl.-Ing. Wolfgang Nejdl
Advisor: Prof. Dr. Avishek Anand

Agenda

Motivation

Problem Statement

Related Work

Our Contribution

Divide, Train and Merge

- Distribution Preserving Sampling
- Merging Approaches

-Experiments Setup

-Performance and Timing

-Conclusion and Works on Going

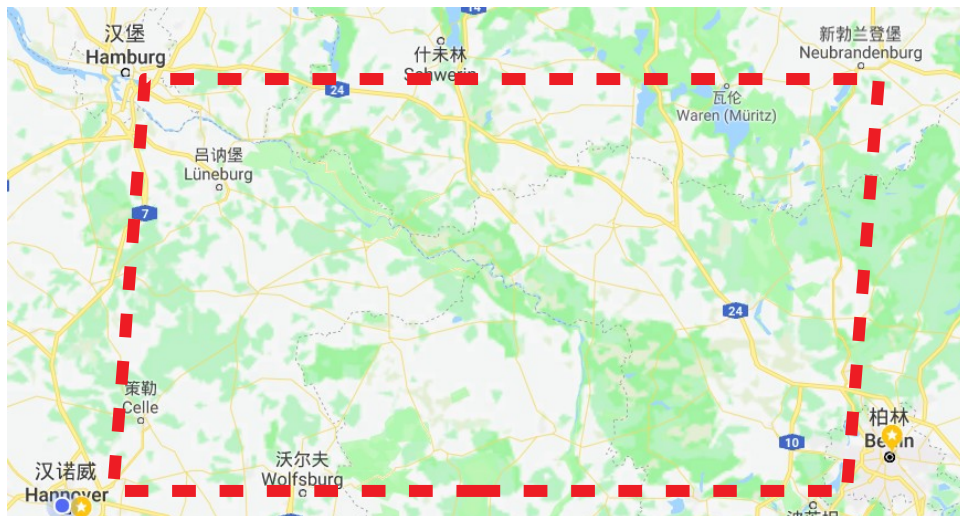
Motivation

1. Hannover – Berlin + Hamburg = ?

Motivation

Hannover – Berlin + Hamburg = Neubrandenburg!

According to the map:



Then what is

King – Man + Woman?

Motivation

A map (the one on a paper or in an App):
maps (mathematical jargon) the city in an area of the surface of
the Earth to a point in E^2 .


If we change “cities” to “words” and the “surface of Earth” to
“Language”

King – Man + Woman


seems solvable!

Motivation


Word Vector Representation



Corpus



Representation model





(Local) Euclidean Vector Space


Motivation

Advantages of Word2Vec


(Mikolov2013)

 $\cos(\theta) \sim \text{similarity}$ 

categories \Rightarrow clusters



analogy \Rightarrow parallelization



Dimension reduction

Problem Statement

Current Challenges:

larger corpus: more **expressive**, longer **time**

On 14GB English Wikidump: \approx **36 hours !**

(dim: 500 Vocabulary: 300k)

Vanilla Word2Vec isn't **scalable**

Related Work

Word2Vec uses SGD to optimize its objective

One solution: make SGD faster:

Optimizing SGD schema (online, synchronized) :

HogWild!^(Recht2011), BLAS-3 SGNS^(Ji2016), cache optimized Hierarchical Softmax^(Eickhoff2016),
column-wise distribution ^(Ordentlich2016)

Adaptive SGD learning rate:

SGD-Momentum^(Rumelhart1986), AdaGrad, RMSProp


Heterogeneous combination:

Combine vectors from Word2Vec, WordNet, GloVe, ConceptNet etc.


Our Contribution

Find a

- scalable
- hardware independent
- asynchronous
- distributional paradigm,
for training word vector representation that
is comparable with the baseline




Divide, Train and Merge



Divide, Train and Merge


Research Questions:



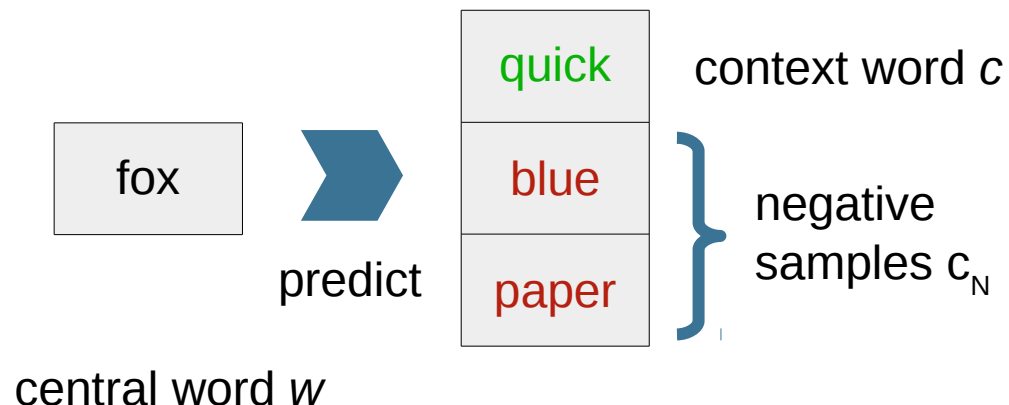
1) How do we sample
from the original corpus?

2) How do we
combine sub-models?

Distribution Preservation Sampling



Skip Gram



Negative Sampling

Distribution Preservation Sampling

Criterion of division: approach to the original term distribution

The local optimum^(Levy2014):


$$\vec{w} \cdot \vec{c} \propto \log \frac{p(w, c)}{p(w) p(c)} \quad \begin{array}{l} \text{Point-wise} \\ \text{Mutual} \\ \text{Information} \end{array}$$

Depends on **distribution of Words/skip-grams** only!


Distribution Preserving Sampling

Distribution preservation through uniform reservoir sampling

KL-divergence of words distribution
from sub-corpus to original corpus




KL-divergence of skip grams distribution
from sub-corpus to original corpus



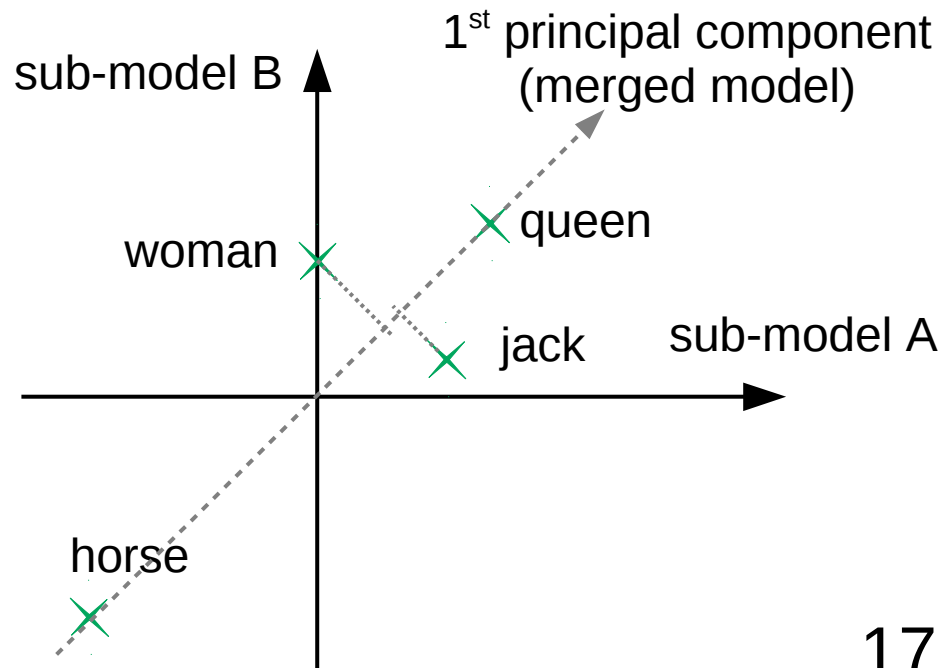
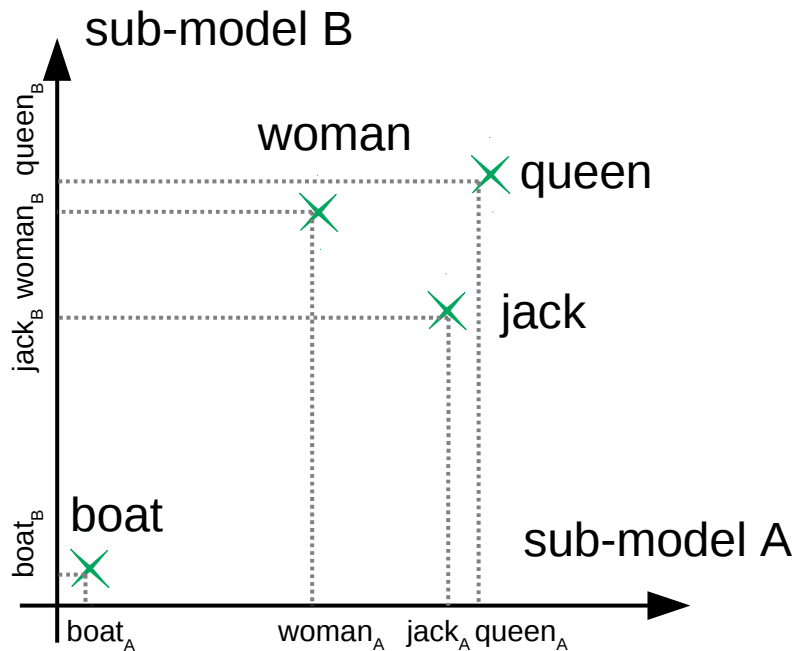
Merging Approaches

Problem so-far:




Merging Approaches

1st Approach: Concatenation + Principle Component Analysis



Merging Approaches

2nd Approach: Manifold Alignment (Boucher2015) + Vector Averaging



[2]

Low-rank reconstruction of X

$$X_{\text{final}}, Y_{\text{final}} = (1-\mu)(\text{reconstructions}) + \mu(\text{correlation})$$

Low-rank reconstruction of Y

$$M_{\text{merged}} = (X_{\text{final}} + Y_{\text{final}}) / 2$$

[2]:

Ham J, Lee D D, Saul L K.
Semisupervised alignment of manifolds
AISTATS. 2005: 120-127.

Experiments Setup

Data Set:

- *Original corpus*: English Wikipedia dump file 2016 (14G)
- *Word2Vec implementation*: gensim^[1]
- *Cleaner*: Wikipedia LineSentences cleaner in Gensim
- *Sample/Partition size*: 1%, 10% of the original corpus

Baseline Model^(Levy2015):

- *Vocabulary*: 300k words
- *Dimension (size of hidden layer)*: 500
- *Training time*: 36 hours

Benchmarks^(Jastrzebski):

- *Word Similarity*: MEN, RW, RG65, WS353, SimLex999, MTurk
- *Analogy*: Google, SemEval 2012 #2, MSR
- *Categorization*: AP, BLESS, Battig

[1]: <https://radimrehurek.com/gensim/>

Performance and Timing

Effect of Dividing (sampling/partitioning):

| Approach | AP | Battig | MEN | RG65 | RW | WS353 | Google | SemEval2012_2 |
|----------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|---------------|
| Partition(1/10,10) | 0.582 | 0.420 | 0.719 | 0.725 | 0.232 | 0.596 | 0.615 | 0.162 |
| Random(1/10,10) | 0.604 | 0.435 | 0.740 | 0.757 | 0.209 | 0.637 | 0.654 | 0.188 |
| Partition(1/100,100) | 0.505 | 0.358 | 0.622 | 0.723 | 0.203 | 0.516 | 0.467 | 0.129 |
| Random(1/100,100) | 0.517 | 0.381 | 0.644 | 0.746 | 0.257 | 0.542 | 0.487 | 0.138 |
| Baseline | 0.595 | 0.434 | 0.736 | 0.757 | 0.299 | 0.611 | 0.661 | 0.181 |

Table 2: Comparison between Random(. , .) and Partition sampling approaches at different granularities. PCA is used as the merging method.

Conclusions:

Sampling performs overall better than partitioning

Larger division size makes the performance better, but less scalability

Surprisingly, some of our approaches even outperform the **baseline**!

Performance and Timing

Effect of Merging (concatenation/PCA/LRA):

| Approach | AP | Battig | MEN | RG65 | RW | WS353 | Google | SemEval2012_2 |
|----------------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|---------------|
| Random(1/10,10) + Concat | 0.595 | 0.429 | 0.742 | 0.765 | 0.268 | 0.611 | 0.645 | 0.187 |
| Random(1/10,10) + LRA | 0.560 | 0.327 | 0.685 | 0.785 | 0.264 | 0.606 | 0.661 | 0.158 |
| Random(1/10,10) + PCA | 0.604 | 0.435 | 0.740 | 0.757 | 0.209 | 0.637 | 0.654 | 0.188 |
| Random(1/100,100) + Concat | 0.525 | 0.377 | 0.643 | 0.746 | 0.257 | 0.542 | 0.487 | 0.138 |
| Random(1/100,100) + LRA | – | – | – | – | – | – | – | – |
| Random(1/100,100) + PCA | 0.517 | 0.381 | 0.644 | 0.746 | 0.257 | 0.542 | 0.487 | 0.138 |
| Baseline | 0.595 | 0.434 | 0.736 | 0.757 | 0.299 | 0.611 | 0.661 | 0.181 |

Table 3: Evaluation results for merging

Conclusions:


PCA performs overall the best

LRA and Vector average is the slowest

To see the information loss in the PCA, we added the concatenation as a single line, it performs surprisingly robust, despite its higher dimensionality

Performance and Timing

Effect of sample count (**10%** of original corpus):




Conclusion:

Larger count of sub-models doesn't help.
8~10 for **10%** approach, **20~30** for **1%** approach is enough.
After that it goes into flat road.

Performance and Timing

Effect of sample count (1% of original corpus):



Conclusion:

Larger count of sub-models doesn't help.


8~10 for 10% approach, 20~30 for 1% approach is enough.

After that it goes into flat road.

Performance and Timing

Timing:

Time of all approaches (the lower the better)



Conclusions:

m-part division leads to m-folds speed up
Total training time: from **36 hours** @ **baseline** to
3.6 hours @ **10%** and
< 1 hour @ **1%** approach
Time of merging is in minutes

Conclusion and on Going Works

Conclusions:

- Sentence-wise sampling to divide, PCA to merge
- M-times speed up by M-fold dividing, sometimes even outperform the baseline.

Ongoing work:

- dealing with missing-word problem between sub-models

Thank you for your attention

Questions and comments are highly appreciated

Zhang, Zijian
Hannover
08. 03. 2018

References:

- Mikolov2013: Mikolov T, Yih W, Zweig G.
Linguistic regularities in continuous space word representations
Proceedings of the 2013 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies. 2013: 746-751.
- Recht2011: Recht B, Re C, Wright S, et al.
Hogwild: A lock-free approach to parallelizing stochastic gradient descent
Advances in neural information processing systems. 2011: 693-701.
- Ji2016: Ji, Shihao, et al.
Parallelizing word2vec in shared and distributed memory
arXiv preprint arXiv:1604.04661 (2016).
- Eickhoff2016: Vuurens, Jeroen BP, Carsten Eickhoff, and Arjen P. de Vries
Efficient Parallel Learning of Word2Vec
arXiv preprint arXiv:1606.07822 (2016).
- Ordentlich2016: Ordentlich, Erik, et al.
Network-efficient distributed Word2vec training system for large vocabularies
Proceedings of the 25th ACM International on Conference on Information and Knowledge Management. ACM, 2016.
- Rumelhart1986: Rumelhart, David E., Geoffrey E. Hinton, and Ronald J. Williams
Learning representations by back-propagating errors
nature 323.6088 (1986): 533.

References:

- Levy2014: Levy O, Goldberg Y.
Neural word embedding as implicit matrix factorization
Advances in neural information processing systems. 2014: 2177-2185.
- Korenius2007: Korenius T, Laurikkala J, Juhola M.
On principal component analysis, cosine and Euclidean measures in information retrieval
Information Sciences, 2007, 177(22): 4893-4905.
- Boucher2015: Boucher, Thomas et al.
Aligning Mixed Manifolds
AAAI. 2015
- Levy2015: Levy, Omer, Yoav Goldberg, and Ido Dagan
Improving distributional similarity with lessons learned from word embeddings
Transactions of the Association for Computational Linguistics 3 (2015): 211-225.
- Jastrzebski2017: Jastrzebski, Stanisław, Damian Leśniak, and Wojciech Marian Czarnecki
How to evaluate word embeddings? On importance of data efficiency and simple supervised tasks
arXiv preprint arXiv:1702.02170 (2017).

(Appendix) Other reasons why W2V is popular

Other reasons why Word2Vec is popular:

Word representations:

used to be processed in computer:

One-hot / hash tokenizing (-lose of semantic information)

LSA: directly SVD on word-doc matrix (-linear model)

Bag-of-Word: meaning defined by set of labels (-low semantic accuracy, supervised)

Words co-occurrence matrix: count of word appears together (-sparse)

Therefore Word2Vec: a predictive model

preserves semantic information

dense representations for larger corpus

non-linear dimension reduction, relationship-preserving and robust

unsupervised way


a neural approach

popular in downstream applications

(Appendix) Some Comments on PCA

(Korenus2007)

Some Comments on PCA [1]



[1]: Korenius T, Laurikkala J, Juhola M.
On principal component analysis, cosine and Euclidean measures in information retrieval
Information Sciences, 2007, 177(22): 4893-4905.